



CS5100 Foundations of Artificial Intelligence

Module 09 Lesson 14

Adversarial Search

Some images and slides are used from CS188 UC Berkeley/AIMA with permission
All materials available at <http://ai.berkeley.edu> / <http://aima.cs.berkeley.edu>

Overview

Minimax
Algorithm

α - β
(Alpha-Beta)
Pruning

Evaluation
functions

Expectiminimax
Algorithm


In this module...

- We consider
 - Multiple agents
 - Competitive (zero-sum), Sequential, and Deterministic (and later Stochastic) environments
 - Using atomic representations
- AI interested in: Deterministic, two-player, perfect info, zero-sum, sequential games
 - E.g. tic-tac-toe (noughts and crosses), chess, checkers, etc.
- Bridge, poker etc. are different – imperfect information, multi-player
 - Some info not visible to all players



Tic-tac-toe or Noughts and Crosses

<https://playtictactoe.org/>



On Checkers



On Chess - Garry Kasparov versus Deep Blue



On Go





Search vs. Games

Search

- Single agent
- Bad path or bad heuristics → slow to get to solution, or no solution

Adversarial Search / Games

- At least 2 agents, may be more
- Bad heuristics/path → (big) loss

Why Games?

Hard to solve!

Can model real-life situations

Rules are clear, and the world is bounded (somewhat!)


- Branching factor for chess = 35,
depth of game ~ 100 (50 moves by each player)
So $35^{100} \sim 10^{154}$ nodes
- Need to make *some* decision, even if not able to calculate optimal decision
- Penalty is severe
- Research into decision-making, and how to make the best use of time
 - Tradeoff: cost to *compute* solution vs. cost of solution

Zero-Sum Games



Zero-Sum Games

- Agents have opposite utilities (values on outcomes) – win for one is a loss for the other
- Think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition



General Games

- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition etc.: all possible

Different Types of Games

Dimensions:


- **Deterministic** or stochastic?
- Single agent or **multi-agent**?
One, two, or more players?
- **Perfect information** (can you see the state)?
- **Sequential** vs. Episodic?
- **Zero sum**?

AI interest in: Deterministic, two-player, perfect info, zero-sum, sequential games

- Wanted: a **strategy (policy)** which recommends a move from each state
- Most games like tic-tac-toe, chess, checkers, etc. fall into this category
- Bridge, poker etc. are different – imperfect information, multi-player
 - Some info not visible to all players
- Physical games (hockey, baseball, tennis) more complicated, lots of actions, imprecise rules (why else would you need a ref?)



Game Playing State-of-the-Art




1994: Chinook beat
Marion Tinsley. 2007:
Checkers is solved!

1997: Deep Blue
beats Garry Kasparov

2016-2017-2019
AlphaGo, AlphaZero,
MuZero ...

Adversarial Search




Deterministic Games


Many possible formalizations, one is:

- States S , Initial State s_0
- Players(s): $P=\{2\dots N\}$ (usually take turns)
- Actions(s): A (may depend on player / state)
- Results(s, a) or Transition Function: $S \times A \rightarrow S$
- Terminal-Test(s): $S \rightarrow \{t, f\}$
- Utility function: $S \times P \rightarrow R$

Solution for a player is a **policy**: $S \rightarrow A$




Single-Agent Search Trees




Further levels

Value of a State

Value of a state:
The best
achievable
outcome (utility)
from that state



Adversarial Game Trees



Minimax Values


Agent's Point of View (here: Pacman's POV)

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$

Tic-Tac-Toe Game Tree: Blue's Perspective



MAX (X)



MIN (O)





MAX (X)



MIN (O)

TERMINAL


Utility



Exercise

Play Tic-Tac-Toe

- Consider this TTT game, after 3 moves by each player
- Assume it is X's turn
- Assume both players are rational
- What should X play?
 - Start with what X could play.
 - Then, what might O respond with?
 - How will it play out...
 - How can you decide from the tree of possibilities?




Approach to Solution

1. Play out the game
2. Label bottom-most states (boards) with +1, 0, -1
 - a. +1 is a win (for the current player)
 - b. 0 is a draw
 - c. -1 is a loss
3. Then label boards one level up, keeping in mind whose move it is
4. Move up level by level

Minimax algorithm, implemented as DFS

Adversarial Search (Minimax)

- Deterministic, zero-sum games:
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- Minimax search:
 - A state-space search tree
 - Players alternate turns
 - Compute each node's **minimax value**: the best achievable utility against a **rational** (optimal) adversary



Minimax Implementation

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

def max-value(state):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

def min-value(state):


initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$


return v

Minimax Example



Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$







Minimax for Chess?

- Does it make sense to use minimax for chess?
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But: do we need to explore the whole tree?

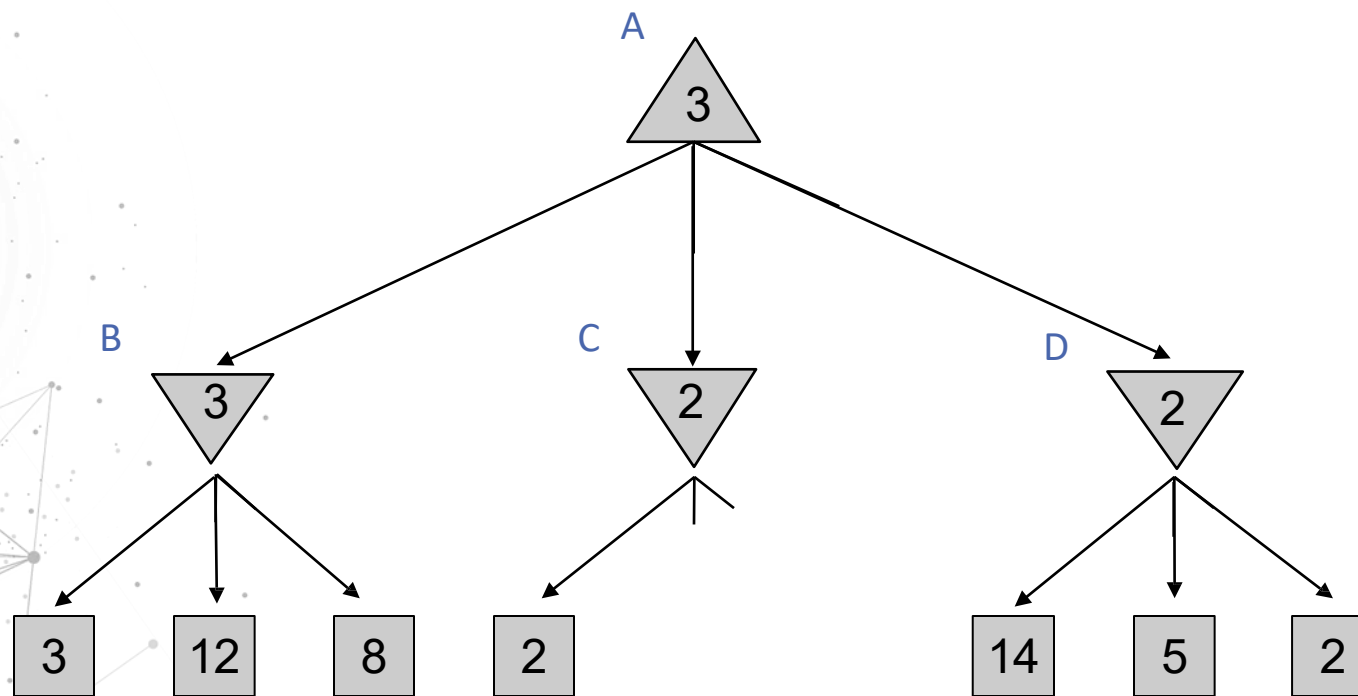
Game Tree Pruning



Minimax Example

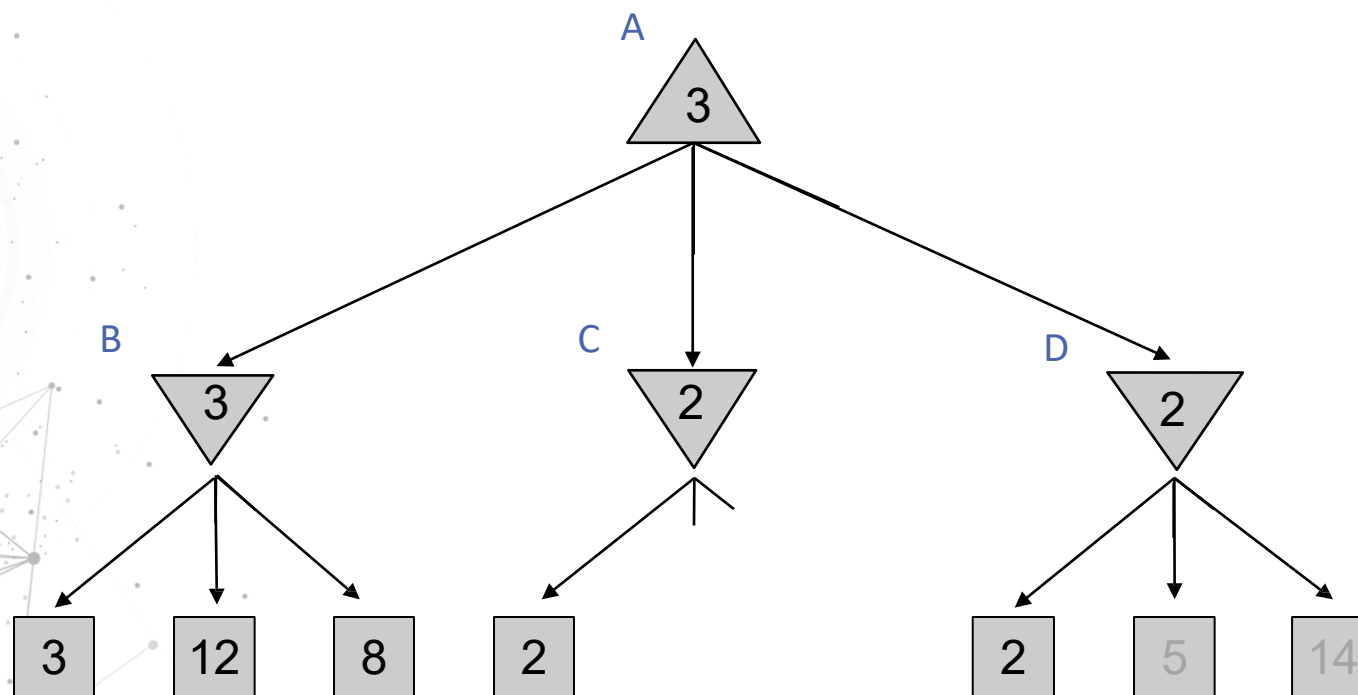


Alpha-Beta Pruning




At each node:
Alpha = at least, Beta = at most

Alpha-Beta Pruning: Reordered subtree



Minimax Algorithm with α - β Pruning




Payoff for Max

Berkeley Video


Alpha-Beta Example

Alpha = best already explored option
along path to the root for maximizer
Beta = best already explored option along
path to the root for minimizer




Alpha-Beta Trials ... 1

Left to Right




Alpha-Beta Trials ... 2

Right to Left




Alpha-Beta Trials ... 3

1,2 → 11, 12



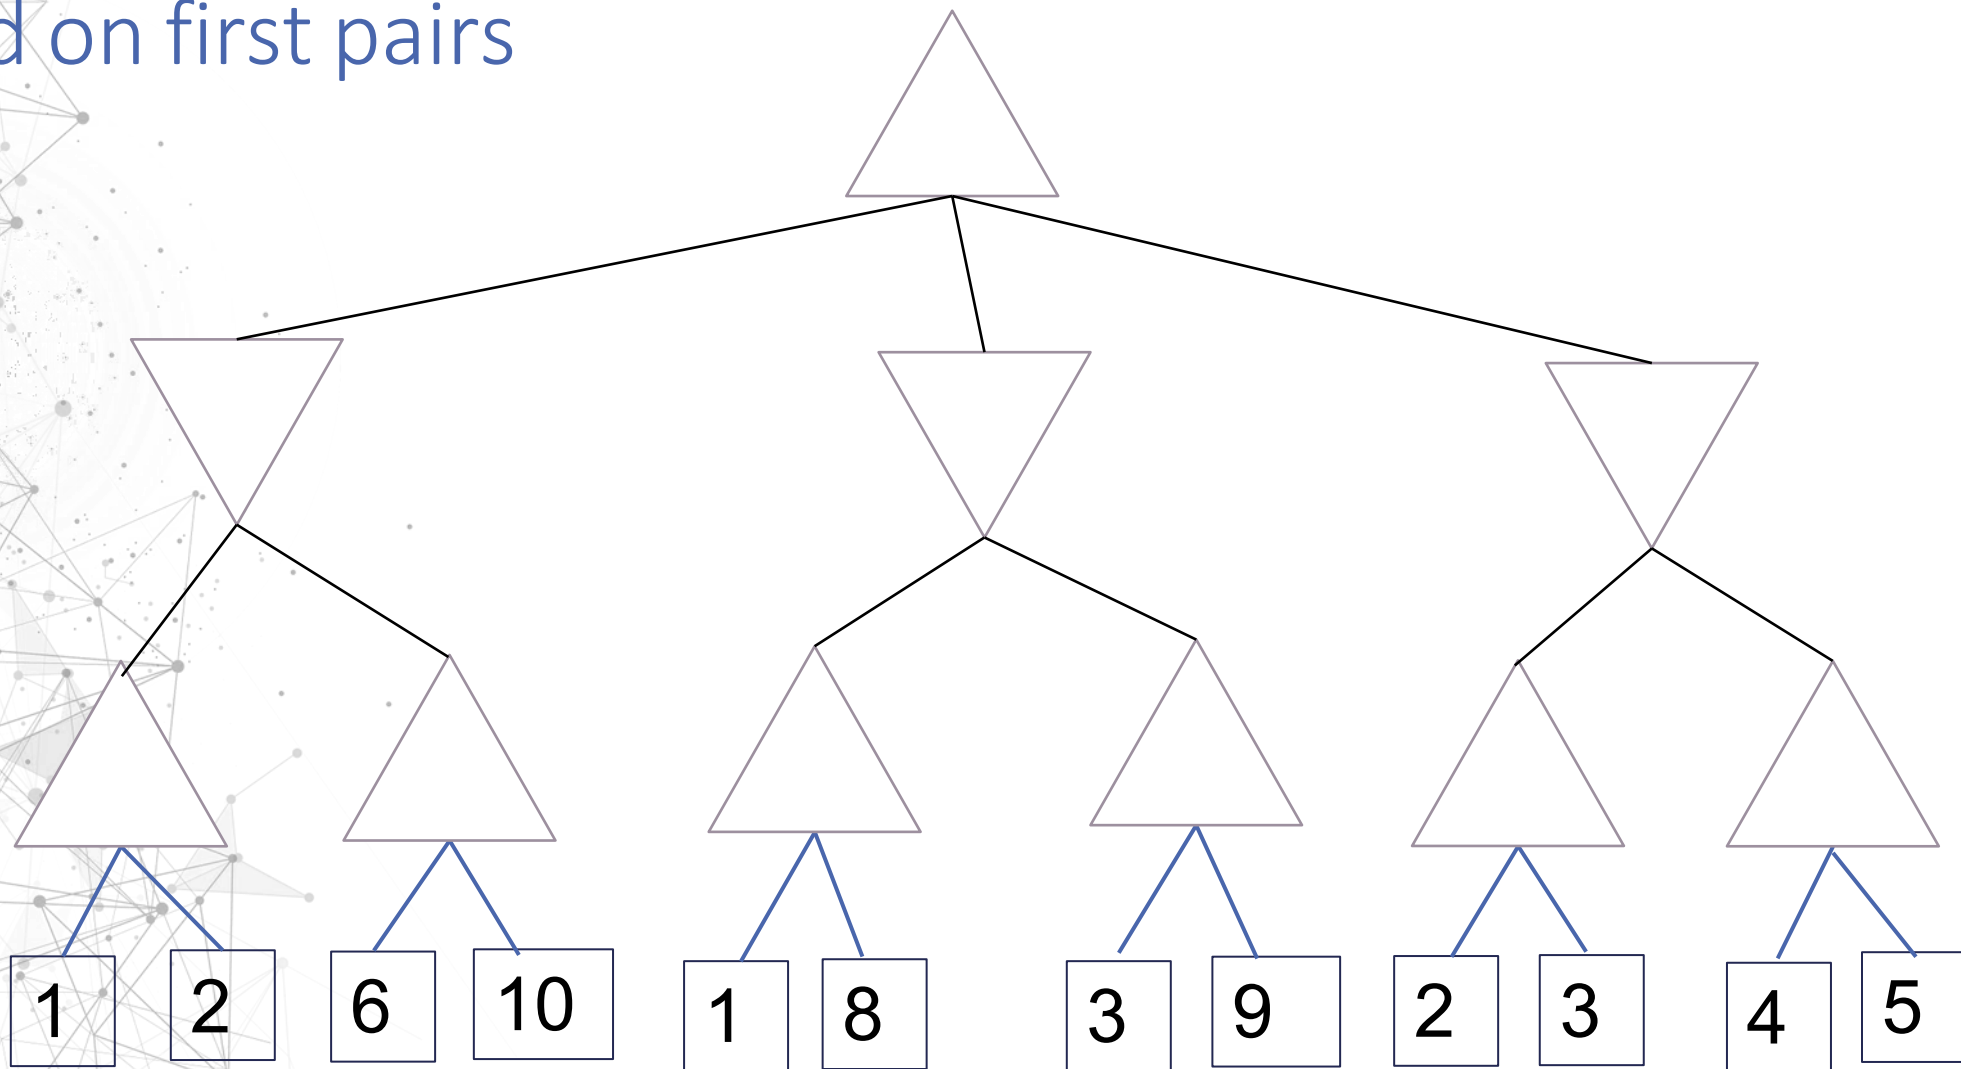
Alpha-Beta Trials ... 4

Sorted (each set of pairs)



Alpha-Beta Trials ... 5

Sorted on first pairs



Alpha-Beta Pruning

General configuration (MIN version)

- We're computing the MIN-VALUE at some node n
- We're looping over n 's children
- n 's estimate of the children's min is dropping
- Who cares about n 's value? MAX
- Let a be the best value that MAX can get at any choice point along the current path from the root
- If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)

MAX version is symmetric


MAX

MIN

⋮

MAX

MIN



Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

Alpha-Beta Pruning Properties

1. This pruning has **no effect** on minimax value computed for the root!


- Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - So the most naïve version won't let you do action selection

2. Good child ordering improves effectiveness of pruning

With “perfect ordering”:


- Time complexity drops to $O(b^{m/2})$
- Doubles solvable depth!
- Full search of, e.g. chess, is still not feasible/sensible...

➔ Simple example of **metareasoning** (computing about what to compute)






A.L. Samuel



A.L. Samuel, 1959: *Some Studies in Machine Learning Using the Game of Checkers* , IBM Jnl, Vol. 3, pp 211-229.

Imperfect Real-Time Decisions & Evaluation Functions



Can't always search all the way to the leaf nodes

You can't just hold up a game till you explore the depths of the tree.

In such cases we may have to make imperfect decisions in real-time.




Basic Idea

Change minimax or alpha-beta

1. Replace utility function with a heuristic evaluation function EVAL which *estimates* position's utility
2. Replace terminal test by a cutoff test which decides when to apply EVAL, e.g., when depth = some depth d , or if terminal state

Resource Limits

- **Problem:** In real games, cannot search to leaves!
- **Solution:** Depth-limited search
 - Instead, search only to a limited depth in the tree
 - At that depth, replace terminal utilities with an evaluation function for non-terminal positions
- Example:
 - Suppose we have 100 seconds, and we can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference





Evaluation Functions

Evaluation functions score non-terminals in depth-limited search



Black to move
White slightly better



White to move
Black winning



EVAL function properties

Ideal function: returns the actual minimax value of the position

In practice, the function:

- Must emulate real utility in ordering terminal states
- Must be fast to compute
- For non-terminal states, must be strongly correlated with chances of winning



EVAL function

In practice: typically, a weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$



e.g. $f_1(s) = (\text{\#white pawns} - \text{\#black pawns}) + (\dots) \dots$

or (difference in piece-count) + (...)

etc.

Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation



Stochastic Games: Dealing with chance



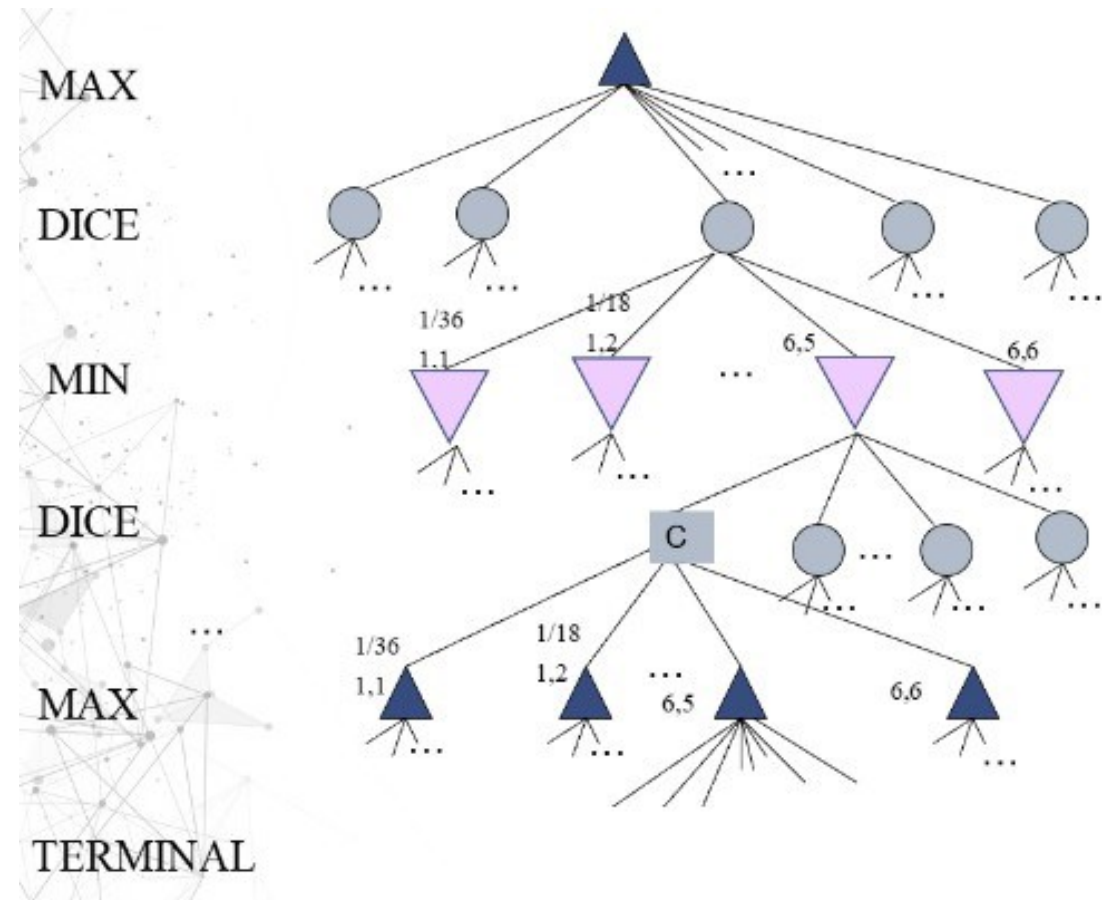
Expectiminimax .. 1

- Generalization of Minimax for games involving chance
- Includes chance nodes between MAX and MIN nodes
- MAX and MIN nodes determined as earlier
- Chance nodes evaluated as “expected value” (weighted average over all possible dice rolls or chance events)



Expectiminimax .. 2

- Includes chance (DICE) nodes between MAX and MIN nodes
- MAX and MIN nodes determined as earlier
- Chance nodes evaluated as “expected value” (weighted average over all possible dice rolls or chance events)
- Branches to chance nodes labeled with probability
 - $6 * 6 = 36$ combinations of dice values, but only 21 distinct
 - 6 doubles at probability $1/36$ each,
 - 15 other combinations at $1/18$ each



Backgammon Game Tree

MAX


DICE

MIN


DICE

MAX

TERMINAL



Computing “Expected Value”



$$2 * 0.9 + 3 * 0.1 = 1.8 + 0.3 = 2.1$$

$$1 * 0.9 + 4 * 0.1 = 0.9 + 0.4 = 1.3$$



Expectiminimax

Expectiminimax(n) =

Utility(n)

for n, a terminal state

$\max_{s \in Succ(n)} \text{expectiminimax}(s)$

for n, a Max node

$\min_{s \in Succ(n)} \text{expectiminimax}(s)$



for n, a Min node

$\sum_{s \in Succ(n)} P(s) * \text{expectiminimax}(s)$

for n, a chance node

Variations of Expectiminimax possible.

Evaluation functions for Expectiminimax



Very different behavior with values at different scales,
emphasizing high payoff

Evaluation functions for games of chance

- Again, cut off search at some level and apply evaluation function to each leaf node
- Dice rolls make things different
- Minimax b^m time
- With chance nodes, need to consider all possible dice rolls sequences, $b^m * N^m$ where $N = \#$ of distinct dice



Card Games

- Bridge, Poker etc. stochastic and partially observable
- Too many possible deals to consider
- Use Monte Carlo approximation, take random sample of N deals, compute best outcome
- Bidding adds even more complexity



