

SQL: Part 1

DML, Relational Algebra

Lecture 3



Relational Algebra

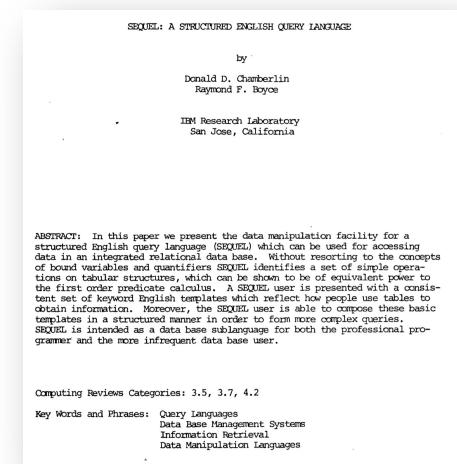
- The basic set of operations for the relational model
 - Note that the relational model assumes sets, so some of the database operations will not map
- Allows the user to formally express a retrieval over one or more relations, as a **relational algebra expression**
 - Results in a new relation, which could itself be queried (i.e. **composable**)
- Why is RA important?
 - Formal basis for SQL
 - Used in query optimization
 - Common vocabulary in data querying technology
 - Sometimes easier to understand the flow of complex SQL



In the Beginning...

Chamberlin, Donald D., and Raymond F. Boyce. "SEQUEL: A structured English query language." *Proceedings of the 1974 ACM SIGFIDET (now SIGMOD) workshop on Data description, access and control.* ACM, 1974.

"In this paper we present the data manipulation facility for a structured English query language (SEQUEL) which can be used for accessing data in an integrated relational data base. Without resorting to the concepts of bound variables and quantifiers SEQUEL identifies a set of simple operations on tabular structures, which can be shown to be of equivalent power to the first order predicate calculus. A SEQUEL user is presented with a consistent set of keyword English templates which reflect how people use tables to obtain information. Moreover, the SEQUEL user is able to compose these basic templates in a structured manner in order to form more complex queries. SEQUEL is intended as a data base sublanguage for both the professional programmer and the more infrequent data base user."



SQL: Structured Query Language

- Declarative: says *what*, not *how*
 - For the most part
- Originally based on relational model/calculus
 - Now industry standards: SQL-86, SQL-92, SQL:1999 (-2023)
 - Various degrees of adoption
- Capabilities
 - Data Definition (DDL): schema structure
 - Data Manipulation (DML): add/update/delete
 - Transaction Management: begin/commit/rollback
 - Data Control: grant/revoke
 - Query
 - Configuration
- ...



Selection

- Our first operation will be to select some tuples from a relation
- This corresponds to the **SELECT** relational algebra operator (σ ; sigma)
 - General form: $\sigma_{<\text{condition}>}(\text{Relation})$
- In SQL this corresponds to the **SELECT** statement



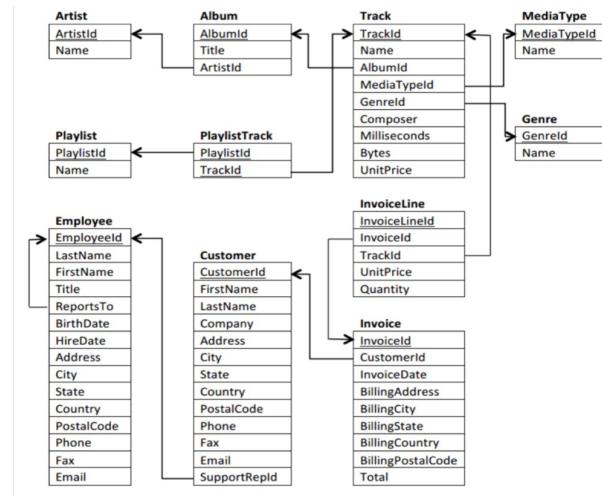
SQL: Simplest Selection

```
SELECT *  
FROM <table name>;
```

Gets all the attributes for all the rows in the specified table. Result set order is arbitrary.



Your First Query!



Get all information about all artists

```
SELECT *
FROM artist;
```

$$\sigma_{true}(artist)$$


Projection/Renaming

- The ability to select a subset of columns from a relation, discarding the rest, is achieved via the PROJECT operator (π ; pi)
 - General form: $\pi_{<\text{attribute list}>}(\text{Relation})$
 - The “attribute list” can include function(s) on existing attributes
- The ability to rename a relation and/or list of attributes is achieved via the RENAME operator (ρ ; rho)
 - General form: $\rho_{<\text{new relation name}>}(<\text{new attribute names}>)(\text{Relation})$
- In SQL these get mapped to the attribute list of the SELECT statement (+ the AS modifier)



SQL: Attribute Control

```
SELECT <attribute list>  
FROM <table name>;
```

Defines the columns of the result set. All rows are returned. Result set order is arbitrary.



Attribute List (1)

- As we saw, to get all attributes in the table, use *

```
SELECT *
FROM employee;
 $\sigma_{\text{true}}(\text{employee})$ 
```

- For a subset, simply list them (comma separated)

```
SELECT FirstName, LastName
FROM employee;
 $\pi_{\text{FirstName}, \text{LastName}}(\sigma_{\text{true}}(\text{employee}))$ 
```

- To rename (or *alias*) an attribute in the result, use **AS**

```
SELECT FirstName AS fname, LastName AS lname
FROM employee;
 $\rho_{(\text{fname}, \text{lname})}(\pi_{\text{FirstName}, \text{LastName}}(\sigma_{\text{true}}(\text{employee})))$ 
```



Attribute List (2)

- In relational algebra, you can optionally show a sequence of steps, giving a name to intermediate relations

$$\rho_{(\text{fname}, \text{lname})}(\pi_{\text{FirstName}, \text{LastName}}(\sigma_{\text{true}}(\text{employee})))$$

vs

$$\text{ALL_E} \leftarrow \sigma_{\text{true}}(\text{employee})$$
$$\text{NAME_E} \leftarrow \pi_{\text{FirstName}, \text{LastName}}(\text{ALL_E})$$
$$\text{RESULT} \leftarrow \rho_{(\text{fname}, \text{lname})}(\text{NAME_E})$$


Attribute List (3)

- In projection, an attribute can be the result of an expression relating existing attributes
 - Available functions depend upon DBMS
 - It is good form to RENAME the result (and makes it easier to access contents via code)

SELECT

```
InvoiceId, InvoiceLineId,  
(UnitPrice*Quantity) AS cost  
FROM invoiceline;
```

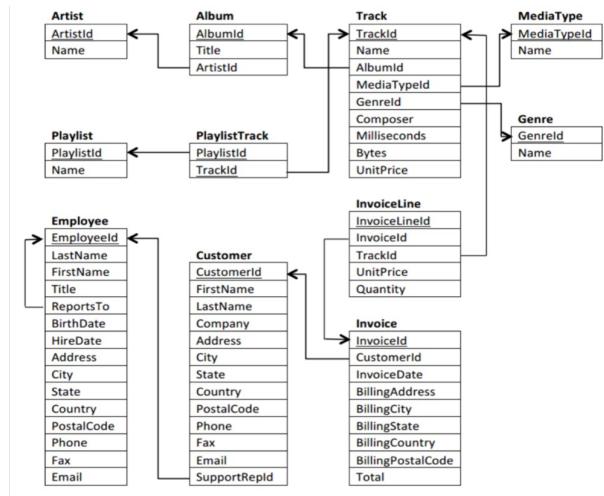
$\text{ALL_ILINES} \leftarrow \sigma_{\text{true}}(\text{invoiceline})$

$\text{ILINE_INFO} \leftarrow \pi_{\text{Invoiceld}, \text{InvoiceLineld}, \text{UnitPrice} * \text{Quantity}}(\text{ALL_ILINES})$

$\text{RESULT} \leftarrow \rho_{(\text{Invoiceld}, \text{InvoiceLineld}, \text{cost})}(\text{ILINE_INFO})$



Basic Queries (1)

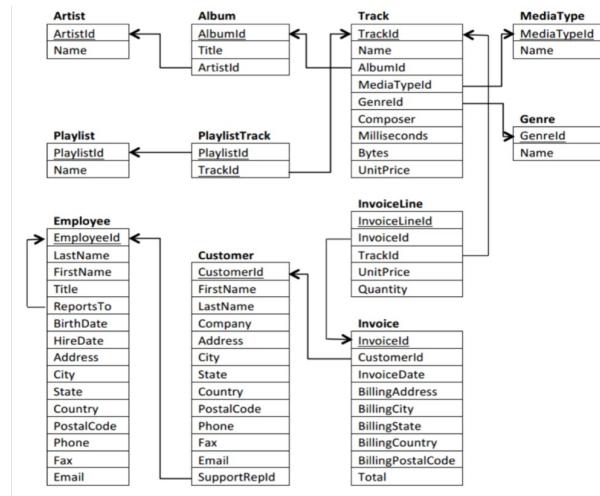


Get all artist names

```
SELECT Name  
FROM artist;
```

$$\pi_{Name}(\sigma_{true}(artist))$$


Basic Queries (2)



Get all employee names (first & last), with their full address info (address, city, state, zip, country)

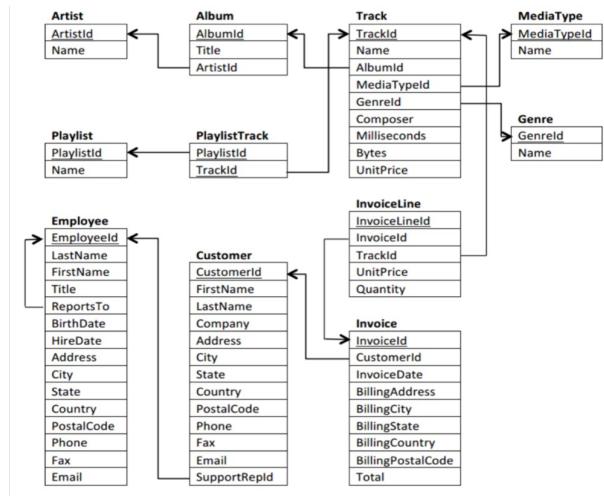
```
SELECT FirstName, LastName, Address, City, State, PostalCode, Country  
FROM employee;
```

$ALL_E \leftarrow \sigma_{true}(employee)$

$RESULT \leftarrow \pi_{FirstName, LastName, Address, City, State, PostalCode, Country}(ALL_E)$



Basic Queries (3)



Get all invoice line(s) with invoice, unit price, quantity

```
SELECT InvoiceId, UnitPrice, Quantity  
FROM invoiceline;
```

$$\pi_{InvoiceId, UnitPrice, Quantity}(\sigma_{true}(invoiceline))$$


Conditional Selection

- Thus far we have included all tuples in a relation
- However, the condition clause of the SELECT operator permits Boolean expressions to restrict included rows
- This corresponds to the **WHERE** clause of the SQL SELECT statement



SQL: Choosing Rows to Include

```
SELECT <attribute list>  
FROM <table name>  
[WHERE <condition list>];
```

Defines the columns of the result set. Only those rows that satisfy the condition(s) are returned. Result set order is arbitrary.



Condition List ~ Boolean Expression

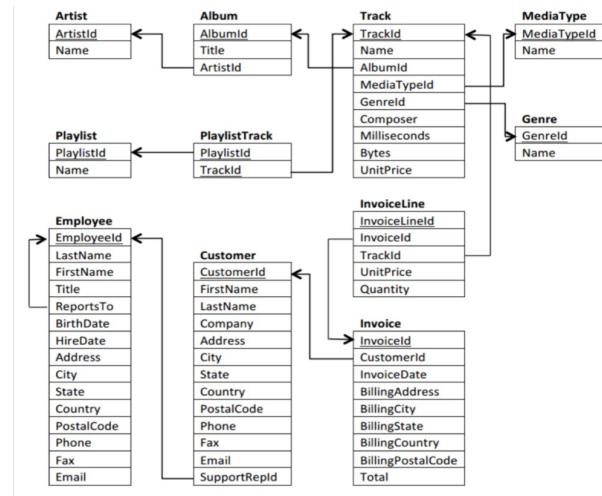
Clauses () separated by **AND/OR**

Operator	Meaning	Example
=	Equal to	<code>InvoiceId = 2</code>
<>	Not equal to	<code>Name <> 'U2'</code>
< or >	Less/Greater than	<code>UnitPrice < 5</code>
<= or >=	Less/Greater than or equal to	<code>UnitPrice >= 0.99</code>
LIKE	Matches pattern	<code>PostalCode LIKE 'T2%</code>
IN	Within a set	<code>City IN ('Calgary', 'Edmonton')</code>
IS or IS NOT	Compare to <code>NULL</code> *	<code>ReportsTo IS NULL</code>
BETWEEN	Inclusive range (esp. dates)	<code>UnitPrice BETWEEN 0.99 AND 1.99</code>

*There are actually is no concept of `NULL` in relational algebra



Conditional Query (1)

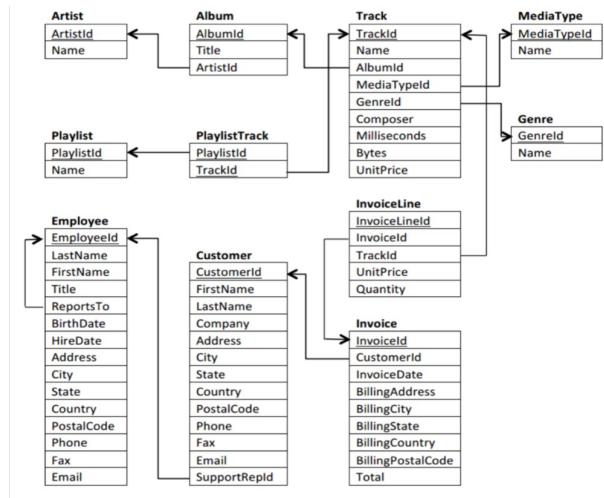


Get the billing country of all invoices totaling more than \$10

```
SELECT BillingCountry
FROM invoice
WHERE Total>10;
```

$$\pi_{BillingCountry}(\sigma_{Total>10}(invoice))$$


Conditional Query (2)

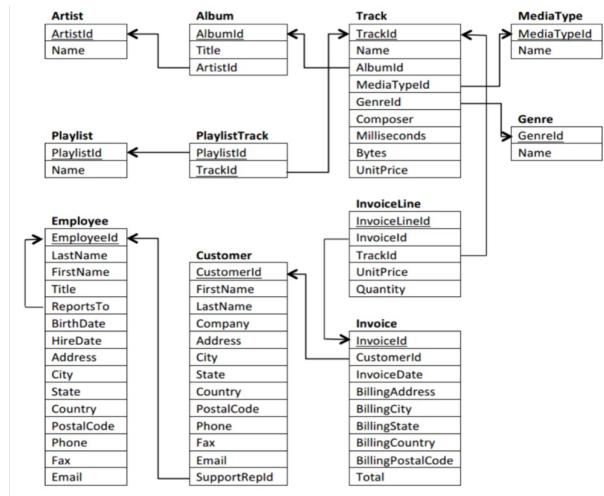


Get all information about tracks whose name contains the word “Rock”

```
SELECT *
FROM track
WHERE Name LIKE '%Rock%';
```

$$\sigma_{Name \text{ } LIKE \text{ } '%Rock%' } (track)$$


Conditional Query (3)



Get the name (first, last) of all non-boss employees in Calgary
(ReportsTo is NULL for the boss).

```

SELECT FirstName, LastName
FROM employee
WHERE ( ReportsTo IS NOT NULL ) AND ( City = 'Calgary' );
    
```

$$\pi_{FirstName, LastName}(\sigma_{ReportsTo \neq EmployeeId \text{ AND } City='Calgary'}(employee))$$

Since RA doesn't have NULL, we could imagine having the Boss report to only themself



Non-Standard Functions

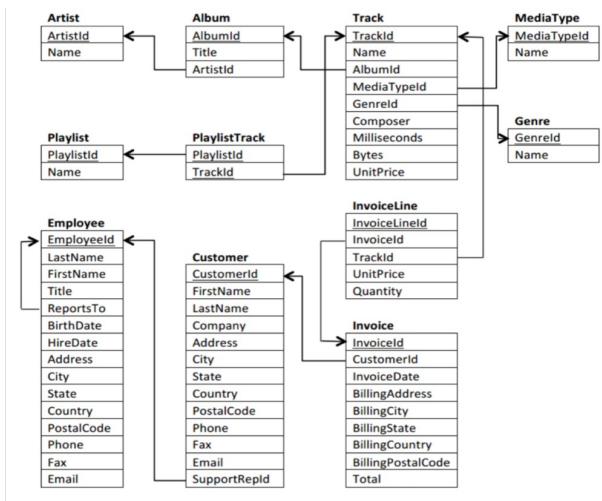
- SQLite
 - <http://sqlite.org/lang.html>
- MariaDB
 - <https://mariadb.com/kb/en/sql-statements/>

Example: Concatenate fields

- SQLite
 - **SELECT (field1 || field2) AS field3**
- MariaDB
 - **SELECT CONCAT(field1, field2) AS field3**



Complex Output Query (SQLite)



Get all German invoices greater than \$1, output the city using the column header “german_city” and “total” prepending \$ to the total

```

1   SELECT BillingCity AS german_city, ('$' || Total) AS total
2   FROM invoice
3   WHERE ( BillingCountry = 'Germany' ) AND ( Total > 1 );
4

```

	german_city	total
1	Stuttgart	\$1.98
2	Berlin	\$1.98
3	Stuttgart	\$13.86
4	Berlin	\$1.98
5	Berlin	\$3.96
6	Berlin	\$13.86
7	Berlin	\$6.94
8	Stuttgart	\$8.91
9	Berlin	\$8.91
10	Frankfurt	\$1.98
11	Frankfurt	\$13.86
12	Frankfurt	\$14.91
13	Stuttgart	\$1.98
14	Stuttgart	\$3.96
15	Berlin	\$1.98
16	Berlin	\$1.98
17	Berlin	\$13.86
18	Stuttgart	\$6.94
19	Berlin	\$3.96
20	Berlin	\$6.94
21	Berlin	\$8.91
22	Frankfurt	\$1.98
23	Frankfurt	\$3.96
24	Frankfurt	\$6.94

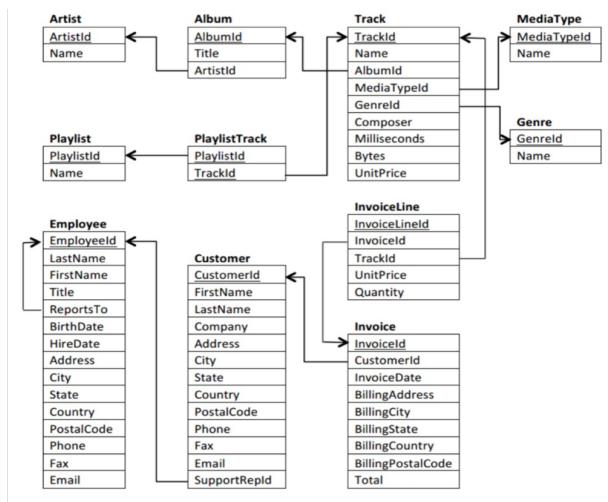
```

SELECT BillingCity AS german_city, ('$' || Total) AS total
FROM invoice
WHERE ( BillingCountry = 'Germany' ) AND ( Total > 1 );

```



Complex Output Query (MariaDB)



Showing rows 0 - 23 (4 total, Query took 0.0000 seconds.)

german_city_total

BillingCity	german_city	total
Stuttgart	\$1.98	
Berlin	\$1.98	
Stuttgart	\$13.86	
Berlin	\$1.98	
Berlin	\$3.96	
Berlin	\$13.86	
Berlin	\$5.94	
Stuttgart	\$8.91	
Berlin	\$8.91	
Frankfurt	\$1.98	
Frankfurt	\$13.86	
Frankfurt	\$14.91	
Stuttgart	\$1.98	
Stuttgart	\$3.96	
Berlin	\$1.98	
Berlin	\$1.98	
Berlin	\$13.86	
Stuttgart	\$5.94	
Berlin	\$3.96	
Berlin	\$5.94	
Berlin	\$8.91	
Frankfurt	\$1.98	
Frankfurt	\$3.96	

Get all German invoices greater than \$1, output the city using the column header “german_city” and “total” prepending \$ to the total

```

SELECT BillingCity AS german_city, CONCAT( '$', Total ) AS total
FROM invoice
WHERE ( BillingCountry = 'Germany' ) AND ( Total > 1 );
    
```

$$G_INV \leftarrow \sigma_{BillingCountry='Germany' \text{ AND } Total>1}(invoice)$$

$$DATA \leftarrow \pi_{BillingCity,CONCAT('$',Total)}(G_INV)$$

$$RES \leftarrow \rho_{(german_city,total)}(DATA)$$

CONCAT is totally non-standard
for relational algebra



SQL: Ordering Output

```
SELECT <attribute list>  
FROM <table name>  
[WHERE <condition list>]  
[ORDER BY <attribute-order list>];
```

Defines the columns of the result set. Only those rows that satisfy the conditions are returned. Result set order is optionally defined.



Relational Algebra Note

- Since the relational model considers relations to be sets (whereas SQL=bags), there is no concept of order
- Some extensions to relational algebra consider that the τ (tau) operator converts the input relation to a bag and outputs an ordered list of tuples
 - General form: $\tau_{<\text{attribute list}>}(\text{Relation})$



SQL: Attribute Order List

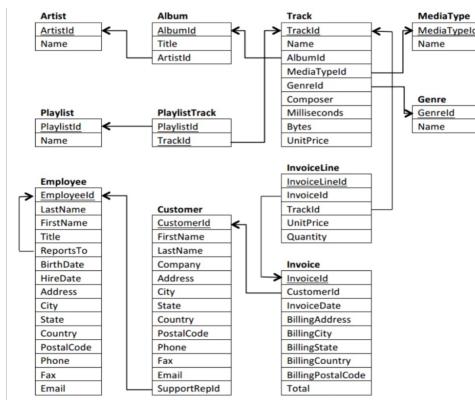
- Comma separated list
- Format: <attribute name> [Order]
 - Order can be **ASC** or **DESC**
 - Default is **ASC**

Example: order all employee information by last name (alphabetical), then first name (alphabetical), then birthdate (youngest first)

```
SELECT *
FROM employee
ORDER BY LastName, FirstName ASC, BirthDate DESC;
```

$$\tau_{LastName, FirstName, BirthDate \text{ DESC}}(\sigma_{true}(employee))$$


Ordering Query



SQL query:

```

1 SELECT *
2 FROM invoice
3 WHERE ( BillingCountry = 'USA' ) AND ( Total >= 10 )
4 ORDER BY Total DESC, BillingState ASC, BillingCity;
5

```

	InvoiceId	CustomerId	InvoiceDate	BillingAddress	BillingCity	BillingState	BillingCountry	BillingPostalCode	Total
1	299	26	2012-08-08 00:00:00	2211 W Berry Street	Fort Worth	TX	USA	76110	23.86
2	201	26	2011-08-29 00:00:00	319 N. Frances Street	Madison	WI	USA	53703	18.86
3	103	24	2010-03-21 00:00:00	1628 B Superior Street	Chicago	IL	USA	60611	18.86
4	397	27	2013-10-13 00:00:00	1033 N Park Ave	Tucson	AZ	USA	85719	18.86
5	26	19	2009-04-14 00:00:00	1 Infinite Loop	Cupertino	CA	USA	95014	13.86
6	124	20	2010-06-22 00:00:00	841 Del Medio Avenue	Mountain View	CA	USA	94040-111	13.86
7	145	16	2010-09-23 00:00:00	1800 Amphitheatre Parkway	Mountain View	CA	USA	94043-1381	13.86
8	320	22	2012-11-06 00:00:00	120 S Orange Ave	Orlando	FL	USA	32801	13.86
9	6	23	2008-01-11 00:00:00	69 Salem Street	Boston	MA	USA	2113	13.86
10	222	21	2011-08-30 00:00:00	801 W 4th Street	Reno	NV	USA	89503	13.86
11	341	18	2013-08-07 00:00:00	827 Broadway	New York	NY	USA	10012-2612	13.86
12	82	28	2009-12-18 00:00:00	302 S 700 E	Salt Lake City	UT	USA	84102	13.86
13	245	17	2011-12-01 00:00:00	1 Microsoft Way	Redmond	WA	USA	98058-8300	13.86
14	311	26	2012-09-28 00:00:00	202 S 700 E	Salt Lake City	UT	USA	84102	11.94
15	298	17	2012-07-31 00:00:00	1 Microsoft Way	Redmond	WA	USA	98058-8300	10.91

Get all invoice info from the USA with greater than or equal to \$10 total, ordered by the total (highest first), and then by state (alphabetical), then by city (alphabetical)

```

SELECT *
FROM invoice
WHERE ( BillingCountry = 'USA' ) AND ( Total >= 10 )
ORDER BY Total DESC, BillingState ASC, BillingCity;

```

$\tau_{Total \text{ DESC}, BillingState, BillingCity}(\sigma_{(BillingCountry='USA') \wedge (Total \geq 10)}(invoice))$



SQL: Set vs. Bag/Multiset

By default, RDBMSs treat results like bags/multisets (i.e. duplicates allowed)

- Use **DISTINCT** to remove duplicates
- For relational algebra, delta: $\delta(\text{Relation})$

```
SELECT [DISTINCT] <attribute list>  
FROM <table name>  
[WHERE <condition list>]  
[ORDER BY <attribute-order list>];
```



Example

```
SELECT BillingState
FROM invoice
WHERE BillingCountry='USA'
ORDER BY BillingState;
```

$$\pi_{BillingState}(\sigma_{BillingCountry='USA'}(\tau_{BillingState}(invoice)))$$

vs.

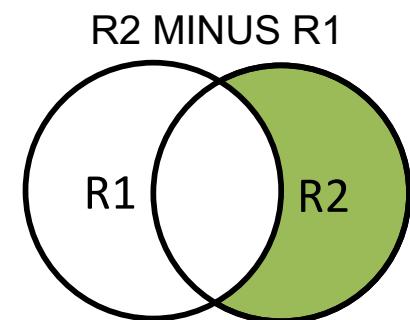
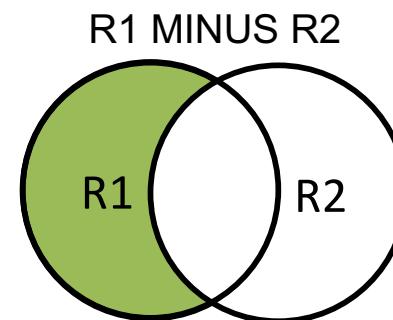
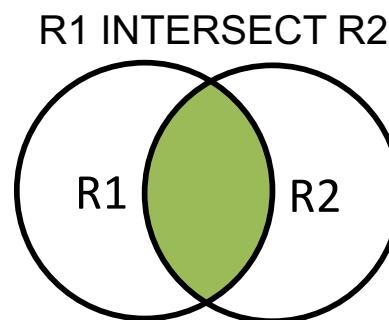
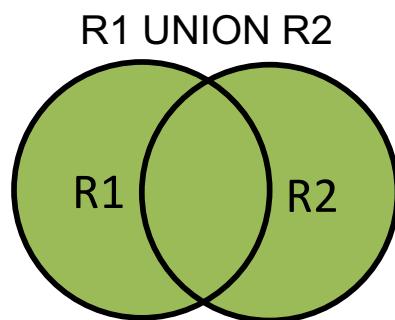
```
SELECT DISTINCT BillingState
FROM invoice
WHERE BillingCountry='USA'
ORDER BY BillingState;
```

$$\delta(\pi_{BillingState}(\sigma_{BillingCountry='USA'}(\tau_{BillingState}(invoice))))$$

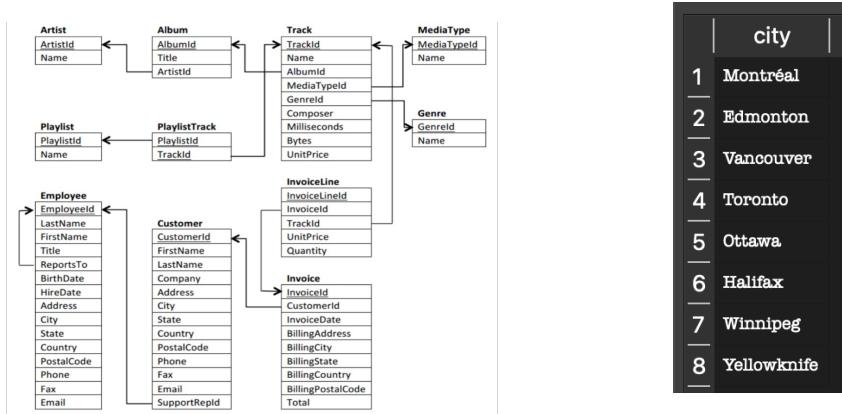

Set Operations

Use **UNION**, **INTERSECT**, **EXCEPT/MINUS** to combine results from queries

- Fields must match exactly in both results
- By default, set handling
 - Use **ALL** after to provide multiset
- Support is spotty here



Combining Queries (1)



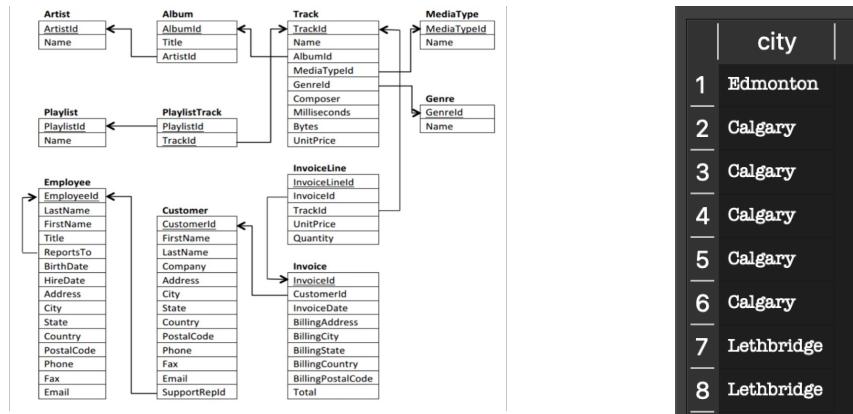
	city
1	Montréal
2	Edmonton
3	Vancouver
4	Toronto
5	Ottawa
6	Halifax
7	Winnipeg
8	Yellowknife

Get all Canadian cities in which customers live
(call result “city”, i.e. lowercase)

```
SELECT City AS city
FROM customer
WHERE Country = 'Canada';
```

$$\rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(customer)))$$


Combining Queries (2)



	city
1	Edmonton
2	Calgary
3	Calgary
4	Calgary
5	Calgary
6	Calgary
7	Lethbridge
8	Lethbridge

Get all Canadian cities in which employees live
(call result “city”, i.e. lowercase)

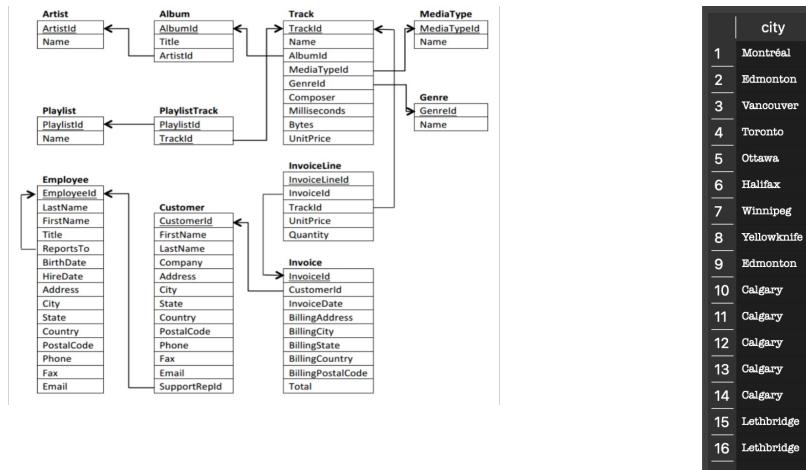
```

SELECT City AS city
FROM employee
WHERE Country = 'Canada';
    
```

$$\rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(employee)))$$



Combining Queries (3)



	city
1	Montréal
2	Edmonton
3	Vancouver
4	Toronto
5	Ottawa
6	Halifax
7	Winnipeg
8	Yellowknife
9	Edmonton
10	Calgary
11	Calgary
12	Calgary
13	Calgary
14	Calgary
15	Lethbridge
16	Lethbridge

Get all Canadian cities in which employees OR customers live (including duplicates)

```
SELECT City AS city FROM customer WHERE Country = 'Canada'
UNION ALL
SELECT City AS city FROM employee WHERE Country = 'Canada';
```

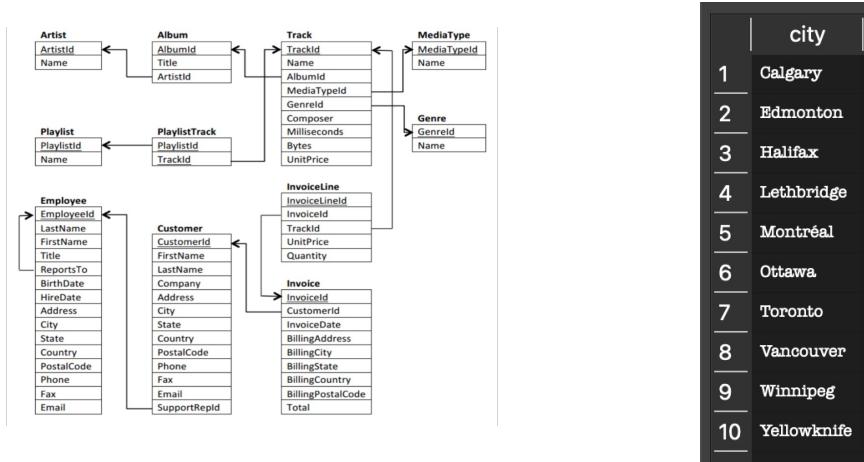
$$R1 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(\tau(customer))))$$

$$R2 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(\tau(employee))))$$

$$RESULT \leftarrow R1 \cup R2$$



Combining Queries (4)



	city
1	Calgary
2	Edmonton
3	Halifax
4	Lethbridge
5	Montréal
6	Ottawa
7	Toronto
8	Vancouver
9	Winnipeg
10	Yellowknife

Get all Canadian cities in which employees OR customers live (excluding duplicates)

```
SELECT City AS city FROM customer WHERE Country = 'Canada'
UNION
SELECT City AS city FROM employee WHERE Country = 'Canada';
```

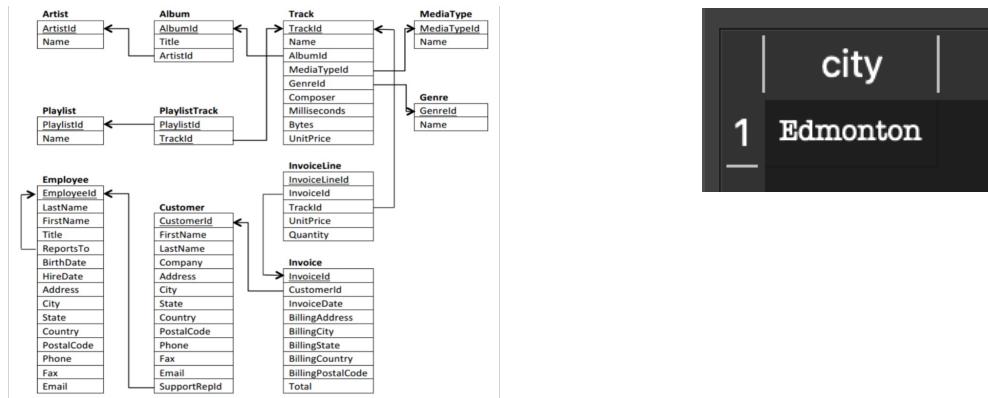
$$R1 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(customer)))$$

$$R2 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(employee)))$$

$$RESULT \leftarrow R1 \cup R2$$



Combining Queries (5)



city	
1	Edmonton

Get all Canadian cities in which employees
AND customers live (excluding duplicates)

```
SELECT City AS city FROM customer WHERE Country = 'Canada'
INTERSECT
SELECT City AS city FROM employee WHERE Country = 'Canada';
```

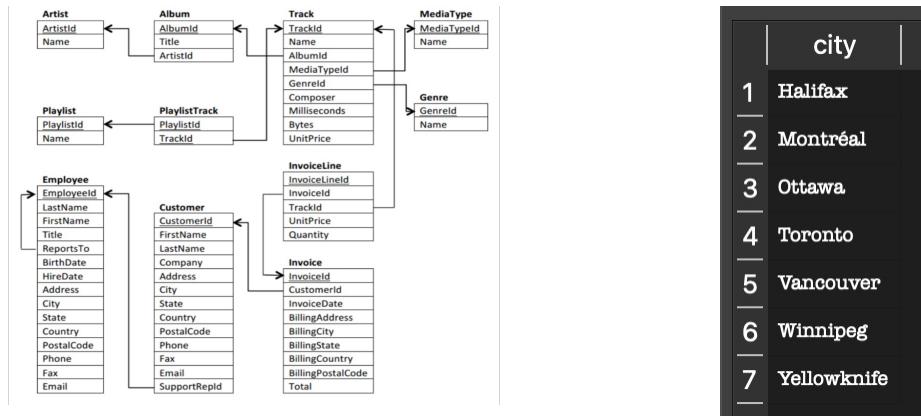
$R1 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(customer)))$

$R2 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(employee)))$

$RESULT \leftarrow R1 \cap R2$



Combining Queries (6)



	city
1	Halifax
2	Montréal
3	Ottawa
4	Toronto
5	Vancouver
6	Winnipeg
7	Yellowknife

All Canadian cities in which customers live BUT employees do not
(excluding duplicates)
[recent MariaDB support]

```
SELECT City AS city FROM customer WHERE Country = 'Canada'
EXCEPT
SELECT City AS city FROM employee WHERE Country = 'Canada';
```

$$R1 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(customer)))$$

$$R2 \leftarrow \rho_{(city)}(\pi_{City}(\sigma_{Country='Canada'}(employee)))$$

$$RESULT \leftarrow R1 - R2$$



Joining Multiple Tables

- SQL supports two methods of **joining** tables, both of which expand the **FROM** clause
 - Basic idea: take Cartesian product of rows, filter
- The first is called a “soft join” and is older and less expressive
 - Not recommended
 - Not covered in detail
- The second uses the **JOIN** keyword and supports more functionality
- Relational algebra: $R_1 \bowtie_{\text{join condition}} R_2$



Intuition: Cartesian Product, Filter (1)

ALPHA

a	b
x	1
y	2
z	3

ALPHA X BETA

Alpha.a	Alpha.b	Beta.c	Beta.d
x	1	x	i
x	1	y	ii
y	2	x	i
y	2	y	ii
z	3	x	i
z	3	y	ii

BETA

c	d
x	i
y	ii



Intuition: Cartesian Product, Filter (2)

ALPHA

a	b
x	1
y	2
z	3

BETA

c	d
x	i
y	ii

ALPHA X BETA | ALPHA.A = BETA.C

Alpha.a	Alpha.b	Beta.c	Beta.d
x	1	x	i
y	2	y	ii
y	2	x	i
y	2	y	ii
z	3	x	i
z	3	y	ii



Simple Join

STUDENT

Name	<u>SSN</u>	Phone	Dorm	Age	GPA	GPA
Ben Bayer	305-61-2435	555-1234	1	19	3.21	3.21
Chung-cha Kim	422-11-2320	555-9876	2	25	3.53	3.25
Barbara Benson	533-69-1238	555-6758	1	19	3.25	

Goal: *find the GPA of students in MATH650*

1. Find all SSN in table Class where Class=MATH650
2. Find all GPA in table Student where SSN=#1

Approach: cross all rows in STUDENT with all rows in CLASS and keep the Student(GPA) of those where STUDENT(SSN)=CLASS(SSN) and CLASS(Class)=MATH650

CLASS

<u>SSN</u>	Class
305-61-2435	COMP355
422-11-2320	COMP355
533-69-1238	MATH650
305-61-2435	MATH650
422-11-2320	BIOL110



Simple Join – JOIN

STUDENT

Name	<u>SSN</u>	Phone	Dorm	Age	GPA	GPA
Ben Bayer	305-61-2435	555-1234	1	19	3.21	3.21
Chung-cha Kim	422-11-2320	555-9876	2	25	3.53	3.25
Barbara Benson	533-69-1238	555-6758	1	19	3.25	

Goal: *find the GPA of students in MATH650*

Approach: cross all rows in STUDENT with all rows in CLASS and keep the GPA of those where STUDENT(SSN)=CLASS(SSN) and CLASS(Class)=MATH650

```
SELECT STUDENT.GPA
FROM STUDENT INNER JOIN CLASS
ON STUDENT.SSN=CLASS.SSN
WHERE CLASS.Class='MATH650';
```

CLASS

<u>SSN</u>	Class
305-61-2435	COMP355
422-11-2320	COMP355
533-69-1238	MATH650
305-61-2435	MATH650
422-11-2320	BIOL110



Simple Join – Soft

STUDENT

Name	<u>SSN</u>	Phone	Dorm	Age	GPA	GPA
Ben Bayer	305-61-2435	555-1234	1	19	3.21	3.21
Chung-cha Kim	422-11-2320	555-9876	2	25	3.53	3.25
Barbara Benson	533-69-1238	555-6758	1	19	3.25	

Goal: *find the GPA of students in MATH650*

Approach: cross all rows in STUDENT with all rows in CLASS and keep the GPA of those where STUDENT(SSN)=CLASS(SSN) and CLASS(Class)=MATH650

```
SELECT STUDENT.GPA
FROM STUDENT, CLASS
WHERE STUDENT.SSN=CLASS.SSN AND
CLASS.Class='MATH650';
```

CLASS

<u>SSN</u>	Class
305-61-2435	COMP355
422-11-2320	COMP355
533-69-1238	MATH650

Soft Joins (older style) intermix
row filtration with
table join conditions



Simple Join – Relational Algebra

STUDENT

Name	<u>SSN</u>	Phone	Dorm	Age	GPA	GPA
Ben Bayer	305-61-2435	555-1234	1	19	3.21	3.21
Chung-cha Kim	422-11-2320	555-9876	2	25	3.53	3.25
Barbara Benson	533-69-1238	555-6758	1	19	3.25	

Goal: *find the GPA of students in MATH650*

Approach: cross all rows in STUDENT with all rows in CLASS and keep the GPA of those where
 $\text{STUDENT(SSN)}=\text{CLASS(SSN)}$ and
 $\text{CLASS(Class)}=\text{MATH650}$

$\text{JOIN} \leftarrow \text{STUDENT} \bowtie_{\text{STUDENT.SSN}=\text{CLASS.SSN}} \text{CLASS}$

$M650 \leftarrow \sigma_{\text{CLASS.Class}='MATH650'}(\text{JOIN})$

$RES \leftarrow \pi_{\text{STUDENT.GPA}}(M650)$

CLASS

<u>SSN</u>	Class
305-61-2435	COMP355
422-11-2320	COMP355
533-69-1238	MATH650
305-61-2435	MATH650
422-11-2320	BIOL110



SQL: Join Syntax

```
SELECT [DISTINCT] <attribute list>  
FROM <table list>  
[WHERE <condition list>]  
[ORDER BY <attribute-order list>];
```

Table List

```
(T1 <join type> T2 [ON <condition list>])  
    <join type> T3 [ON <condition list>]...
```



Join Types

[INNER] JOIN $A \bowtie B$	Row must exist in <u>both</u> tables
LEFT [OUTER] JOIN $A \bowtie B$	Row must <i>at least</i> exist in the table to the left (padded with NULL)
RIGHT [OUTER] JOIN $A \bowtie B$	Row must exist <i>at least</i> in the table to the right (padded with NULL)
FULL OUTER JOIN $A \bowtie B$	Row exists in <u>either</u> table (padded with NULL)



Join Type Example (1)

ALPHA

a	b
x	1
y	2
z	3

SELECT ***FROM Alpha INNER JOIN Beta ON
Alpha.a=Beta.c**
$$\text{Alpha} \bowtie_{\text{Alpha.a}=\text{Beta.c}} \text{Beta}$$
BETA

c	d
w	-
y	ii

Alpha.a	Alpha.b	Beta.c	Beta.d
y	2	y	ii



Join Type Example (2)

ALPHA

a	b
x	1
y	2
z	3

```
SELECT *
FROM Alpha LEFT OUTER JOIN Beta ON
Alpha.a=Beta.c
```

Alpha $\bowtie_{Alpha.a=Beta.c}$ *Beta*

BETA

c	d
w	-
y	ii

Alpha.a	Alpha.b	Beta.c	Beta.d
x	1	NULL	NULL
y	2	y	ii
z	3	NULL	NULL



Join Type Example (3)

ALPHA

a	b
x	1
y	2
z	3

```
SELECT *
FROM Alpha RIGHT OUTER JOIN Beta ON
Alpha.a=Beta.c
```

$\text{Alpha} \bowtie_{\text{Alpha.a}=\text{Beta.c}} \text{Beta}$

BETA

c	d
w	-
y	ii

Alpha.a	Alpha.b	Beta.c	Beta.d
y	2	y	ii
NULL	NULL	w	-



Join Type Example (4)

ALPHA

a	b
x	1
y	2
z	3

```
SELECT *
FROM Alpha FULL OUTER JOIN Beta ON
Alpha.a=Beta.c
```

Alpha $\bowtie_{Alpha.a=Beta.c}$ *Beta*

BETA

c	d
w	-
y	ii

Alpha.a	Alpha.b	Beta.c	Beta.d
x	1	NULL	NULL
y	2	y	ii
z	3	NULL	NULL
NULL	NULL	w	-

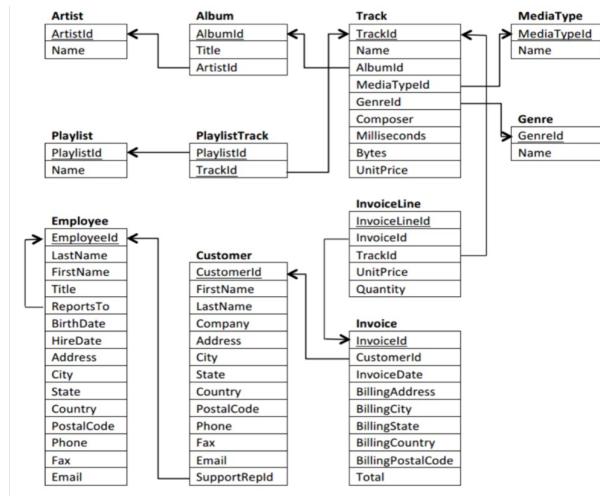


Notes on Joins

- When dealing with multiple tables, it is advised to use full attribute addressing (table.attribute) to avoid confusion
 - Tip: when listing the table name, give it a shortcut
`SELECT * FROM table1 t1`
 $\sigma_{true}(\rho_{t1}(table1))$
- NATURAL ($R_1 * R_2$)
 - Optional shortcut if joining attribute(s) have same name(s) in both tables
- Support/syntax can be spotty
 - Particularly full outer, natural
- When joining, the new set of available attributes (*) is the concatenation of the attributes from *both* tables



Exploring Joins (1)

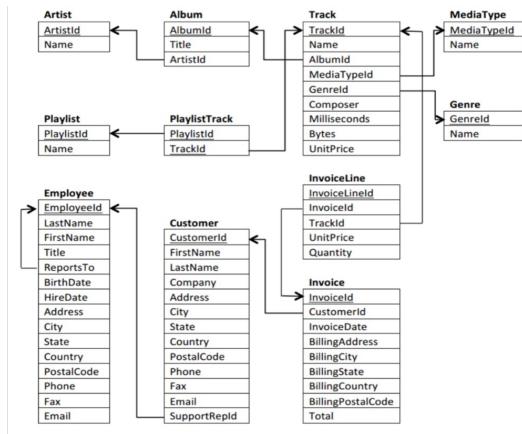


Get the cross product of genres and media types

```
SELECT *
FROM genre INNER JOIN mediatype;
```

$$\sigma_{true}(genre \bowtie mediatype)$$


Exploring Joins (2)



Get all track information, with the appropriate genre name and media type name, for all jazz tracks where Miles Davis helped compose

```

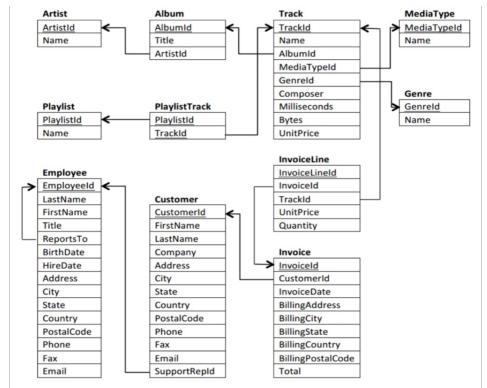
SELECT *
FROM (track t INNER JOIN mediatype mt ON t.MediaTypeId=mt.MediaTypeId)
INNER JOIN genre g ON t.GenreId=g.GenreId
WHERE g.Name='Jazz' AND t.Composer LIKE '%Miles Davis%';
    
```

$$J1 \leftarrow \rho_t(\text{track}) \bowtie_{t.MediaTypeId=mt.MediaTypeId} \rho_{mt}(\text{mediatype})$$

$$J2 \leftarrow J1 \bowtie_{t.GenreId=g.GenreId} \rho_g(\text{genre})$$

$$RES \leftarrow \sigma_{g.Name='Jazz' \text{ AND } t.Composer \text{ LIKE } '%MilesDavis%'}(J2)$$


Advanced Joins (1)



	ArtistId	Name
1	169	Black Eyed Peas
2	11	Black Label Society
3	12	Black Sabbath

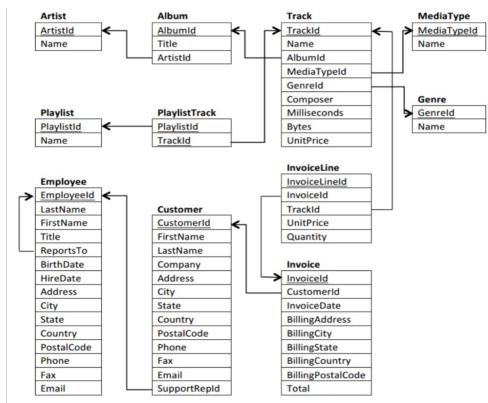
Get all artist information for those whose name begins with ‘Black’, sort by name (alphabetically)

```

SELECT *
FROM artist
WHERE Name LIKE 'Black%'
ORDER BY Name ASC;
    
```

$$\tau_{Name}(\sigma_{Name \text{ } LIKE \text{ } 'Black\%'}(artist))$$


Advanced Joins (2)



	ArtistId	Name	AlbumId	Title	ArtistId
1	11	Black Label Society	14	Alcohol Fueled Brewtality Live! [Disc 1]	11
2	11	Black Label Society	15	Alcohol Fueled Brewtality Live! [Disc 2]	11
3	12	Black Sabbath	16	Black Sabbath	12
4	12	Black Sabbath	17	Black Sabbath Vol. 4 (Remaster)	12

Get all artist AND album information for those artists whose name begins with ‘Black’ (don’t include those without albums), sort by artist name, then album name

```

SELECT *
FROM artist art INNER JOIN album alb ON art.ArtistId=alb.ArtistId
WHERE Name LIKE 'Black%'
ORDER BY art.Name ASC, alb.Title ASC;
    
```

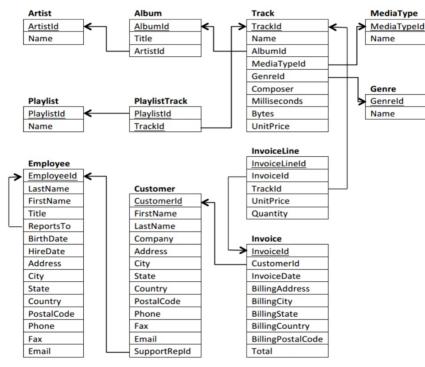
$J \leftarrow \rho_{art}(artist) \bowtie_{art.ArtistId=alb.ArtistId} \rho_{alb}(album)$

$S \leftarrow \sigma_{Name \text{ LIKE } 'Black\%'}(J)$

$RES \leftarrow \tau_{art.Name, alb.Title}(S)$



Advanced Joins (3)



	ArtistId	Name	AlbumId	Title	ArtistId
1	169	Black Eyed Peas	NULL	NULL	NULL
2	11	Black Label Society	14	Alcohol Fueled Brewtality Live! [Disc 1]	11
3	11	Black Label Society	15	Alcohol Fueled Brewtality Live! [Disc 2]	11
4	12	Black Sabbath	16	Black Sabbath	12
5	12	Black Sabbath	17	Black Sabbath Vol. 4 (Remaster)	12

Get all artist AND album information for those artists whose name begins with ‘Black’ (do include those without albums!), sort by artist name, then album title

```

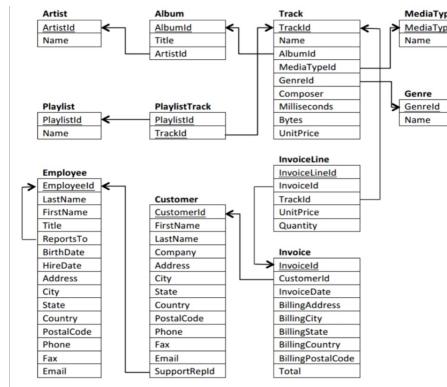
SELECT *
FROM artist art LEFT OUTER JOIN album alb ON art.ArtistId=alb.ArtistId
WHERE Name LIKE 'Black%'
ORDER BY art.Name, alb.Title;
    
```

$$J \leftarrow \rho_{art}(artist) \bowtie_{art.ArtistId=alb.ArtistId} \rho_{alb}(album)$$

$$S \leftarrow \sigma_{Name \text{ LIKE } 'Black\%'}(J)$$

$$RES \leftarrow \tau_{art.Name, alb.Title}(S)$$


Advanced Joins (4)



	ArtistId	Name	AlbumId	Title
1	169	Black Eyed Peas	NULL	NULL
2	11	Black Label Society	14	Alcohol Fueled Brewtality Live! [Disc 1]
3	11	Black Label Society	15	Alcohol Fueled Brewtality Live! [Disc 2]
4	12	Black Sabbath	16	Black Sabbath
5	12	Black Sabbath	17	Black Sabbath Vol. 4 (Remaster)

Get all artist AND album information for those artists whose name begins with ‘Black’ (do include those without albums!), provide only a single correct ArtistId, sort by artist name, then album title

```

SELECT art.ArtistId, art.Name, alb.AlbumId, alb.Title
FROM artist art LEFT OUTER JOIN album alb ON art.ArtistId=alb.ArtistId
WHERE Name LIKE 'Black%'
ORDER BY art.Name, alb.Title;
    
```

$$J \leftarrow \rho_{art}(artist) \bowtie_{art.ArtistId=alb.ArtistId} \rho_{alb}(album)$$

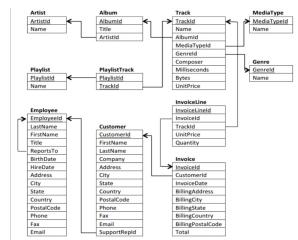
$$S \leftarrow \sigma_{Name \text{ LIKE } 'Black\%'}(J)$$

$$P \leftarrow \pi_{art.ArtistId, art.Name, alb.AlbumId, alb.Title}(S)$$

$$RES \leftarrow \tau_{art.Name, alb.Title}(P)$$



Advanced Joins (5)



TrackId	tName	Composer	UnitPrice	Title	mName	gName	
1	1139	Give Me Novacaine	Green Day	0.99	American Idiot	MPEG audio file	Alternative & Punk

Get track id, track name, composer, unit price, album title, media type name, and genre for the track titled “Give Me Novacaine”

```
SELECT t.TrackId, t.Name AS tName, t.Composer, t.UnitPrice,
       a.Title, m.Name AS mName, g.Name AS gName
  FROM ((track t INNER JOIN album a ON t.AlbumId=a.AlbumId)
        INNER JOIN mediatype m ON t.MediaTypeId=m.MediaTypeId)
        INNER JOIN genre g ON t.GenreId=g.GenreId
 WHERE t.Name='Give Me Novacaine';
```

$$\begin{aligned}
 TA &\leftarrow \rho_t(\text{track}) \bowtie_{t.AlbumId=a.AlbumId} \rho_a(\text{album}) \\
 M &\leftarrow TA \bowtie_{t.MediaTypeId=m.MediaTypeId} \rho_m(\text{mediatype}) \\
 G &\leftarrow M \bowtie_{t.GenreId=g.GenreId} \rho_g(\text{genre}) \\
 S &\leftarrow \sigma_{t.Name='Give Me Novacaine'}(G) \\
 P &\leftarrow \pi_{t.TrackId, t.Name, t.Composer, t.UnitPrice, a.Title, m.Name, g.Name}(S) \\
 RES &\leftarrow \rho_{(TrackId, tName, Composer, UnitPrice, Title, mName, gName)}(P)
 \end{aligned}$$
